

## Assignment II - Valgrind & Pytorch profiler

### 1. Memcheck

Error 1 : Invalid write

```
Invalid write of size 4                                ##error 1 : invalid write
    at 0x1091C0: main (memleak.c:49)
Address 0x4a96068 is 0 bytes after a block of size 40 alloc'd
    at 0x484884F: malloc (vg_replace_malloc.c:393)
    by 0x10919E: main (memleak.c:46)
```

這個錯誤是指程式想要對於 malloc 的記憶體以外的位置進行寫入的動作，  
這個例子是在 malloc 40 bytes 之後，在第 41 byte 的位置進行寫入。

Error 2 : Invalid read

```
Invalid read of size 4                                ##error 2 : invalid read
    at 0x1091ED: main (memleak.c:54)
Address 0x4a96068 is 0 bytes after a block of size 40 alloc'd
    at 0x484884F: malloc (vg_replace_malloc.c:393)
    by 0x10919E: main (memleak.c:46)
```

這個錯誤是指程式想要對於 malloc 的記憶體以外的位置進行讀取的動作，  
這個例子是在 malloc 40 bytes 之後，在第 41 byte 的位置進行讀取。

Error 3 : Conditional jump or move depends on uninitialized

```
==93707== Conditional jump or move depends on uninitialised value(s)
##error 3 : Conditional jump or move depends on uninitialised value(s)
==93707==    at 0x48E0AD6: __vfprintf_internal (vfprintf-internal.c:1516)
==93707==    by 0x48CA79E: printf (printf.c:33)
==93707==    by 0x109214: main (memleak.c:57)
```

這個錯誤是指在條件判斷時使用了未初始化的數值，這個例子是在 printf 裡  
使用未初始化的變數或記憶體，且這個數值影響了條件的判斷。

Error 4 : Use of uninitialized value

```
Use of uninitialised value of size 8                  ##error 4 : Use of uninitialised value
    at 0x48C42EB: _itoa_word (_itoa.c:177)
    by 0x48DFABD: __vfprintf_internal (vfprintf-internal.c:1516)
    by 0x48CA79E: printf (printf.c:33)
    by 0x109214: main (memleak.c:57)
```

這個錯誤是指程式想要讀取或使用操作一個未初始化的數值，這個例子可能  
是要 print 出一個變數，並且在 itoa 時用了一個未初始化的變數。

## Error 5 : Argument 'size' of function malloc has a fishy value

```
==93707== Argument 'size' of function malloc has a fishy (possibly negative) value: -40
##error 5 : Argument 'size' of function malloc has a fishy (possibly negative) value
==93707==    at 0x484884F: malloc (vg_replace_malloc.c:393)
==93707==    by 0x109220: main (memleak.c:61)
```

這個錯誤是指在使用動態分配記憶體 `malloc` 時，給定的大小不是正數。這個例子是在呼叫 `malloc` 函數時給了-40 的大小。

## Error 6 : Invalid free() / delete / delete[] / realloc()

```
==93707== Invalid free() / delete / delete[] / realloc()
##error 6 : Invalid free() / delete / delete[] / realloc()
==93707==    at 0x484B0C4: free (vg_replace_malloc.c:884)
==93707==    by 0x10924A: main (memleak.c:65)
==93707== Address 0x4a964f0 is 0 bytes inside a block of size 40 free'd
==93707==    at 0x484B0C4: free (vg_replace_malloc.c:884)
==93707==    by 0x10923E: main (memleak.c:64)
==93707== Block was alloc'd at
==93707==    at 0x484884F: malloc (vg_replace_malloc.c:393)
==93707==    by 0x10922E: main (memleak.c:63)
```

這個錯誤可能是指以下幾種錯誤，對同一個指標重複 `free`，對一個未分配記憶體的指標 `free`，`free` 一個 local variable。這個例子是對指向一塊大小為 40bytes 的記憶體的 pointer 呼叫了兩次 `free()`。

## 2. Cachegrind

Good:

```
==4418== Cachegrind, a cache and branch-prediction profiler
==4418== Copyright (C) 2002-2017, and GNU GPL'd, by Nicholas Nethercote et al.
==4418== Using Valgrind-3.20.0 and LibVEX; rerun with -h for copyright info
==4418== Command: ./good
==4418== Parent PID: 3209
==4418==
--4418-- warning: L3 cache found, using its data for the LL simulation.
--4418-- warning: specified LL cache: line_size 64  assoc 16  total_size 12,582,912
--4418-- warning: simulated LL cache: line_size 64  assoc 24  total_size 12,582,912
==4418==
==4418== I  refs:      30,158,916
==4418== I1 misses:      1,069
==4418== LLi misses:      1,057
==4418== I1 miss rate:    0.00%
==4418== LLi miss rate:  0.00%
==4418==
==4418== D  refs:      14,053,786 (12,039,693 rd + 2,014,093 wr)
==4418== D1 misses:      127,168 (   64,086 rd +   63,082 wr)
==4418== LLd misses:      64,330 (    1,291 rd +   63,039 wr)
==4418== D1 miss rate:    0.9% (    0.5% +    3.1% )
==4418== LLd miss rate:  0.5% (    0.0% +    3.1% )
==4418==
==4418== LL refs:      128,237 (   65,155 rd +   63,082 wr)
==4418== LL misses:      65,387 (    2,348 rd +   63,039 wr)
==4418== LL miss rate:    0.1% (    0.0% +    3.1% )
```


Bad:

```
==4436== Cachegrind, a cache and branch-prediction profiler
==4436== Copyright (C) 2002-2017, and GNU GPL'd, by Nicholas Nethercote et al.
==4436== Using Valgrind-3.20.0 and LibVEX; rerun with -h for copyright info
==4436== Command: ./bad
==4436== Parent PID: 3209
==4436==
--4436-- warning: L3 cache found, using its data for the LL simulation.
--4436-- warning: specified LL cache: line_size 64  assoc 16  total_size 12,582,912
--4436-- warning: simulated LL cache: line_size 64  assoc 24  total_size 12,582,912
==4436==
==4436== I   refs:      30,158,921
==4436== I1 misses:      1,070
==4436== L1i misses:      1,057
==4436== I1 miss rate:      0.00%
==4436== L1i miss rate:      0.00%
==4436==
==4436== D   refs:      14,053,786 (12,039,693 rd + 2,014,093 wr)
==4436== D1 misses:      2,002,168 ( 1,001,586 rd + 1,000,582 wr)
==4436== L1d misses:      64,330 (   1,291 rd +   63,039 wr)
==4436== D1 miss rate:     14.2% (    8.3% +   49.7% )
==4436== L1d miss rate:    0.5% (    0.0% +    3.1% )
==4436==
==4436== LL refs:         2,003,238 ( 1,002,656 rd + 1,000,582 wr)
==4436== LL misses:         65,387 (    2,348 rd +    63,039 wr)
==4436== LL miss rate:      0.1% (    0.0% +    3.1% )
```

Good 和 Bad 的差別在於 D1 misses，這代表 bad 在 level 1 data cache 有很多的 miss。這有可能是因為 data access pattern 不同的緣故，bad 很可能是以 strided 或 random 的取法，導致 cache miss 一直發生。除此之外，也有可能是 good 的 data locality 比較好。而 D1 misses 也影響到了 LL refs，因為 cache 是從上層取到下層，而前面層數越多 miss，last level 也就有可能會有越多的 reference。

3. Massif

(1)



(2)

n	time(B)	total(B)	useful-heap(B)	extra-heap(B)	stacks(B)
60	225,568	225,568	225,000	568	0
61	229,576	229,576	229,000	576	0
62	231,584	231,584	231,000	584	0
63	235,592	235,592	235,000	592	0
64	239,600	239,600	239,000	600	0
65	239,600	239,600	239,000	600	0


Allocated: 239600 bytes

Used: 239000 bytes

#### 4. Callgrind

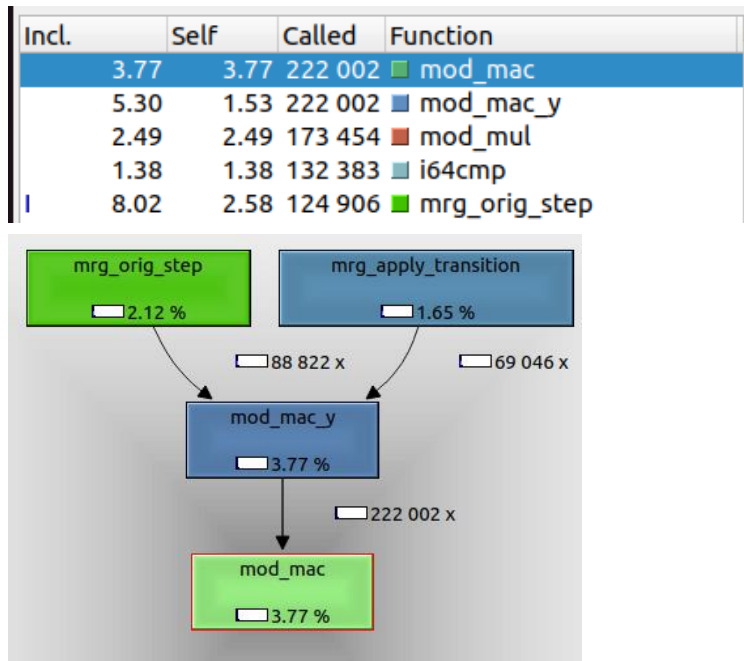
(1)

Incl.	Self	Called	Function
54.50	48.33	64	verify_bfs_tree
11.34	11.33	64	make_bfs_tree
6.16	6.16	64	compute_levels
3.77	3.77	222 002	mod_mac
16.43	3.13	12 288	make_one_edge
10.76	2.80	24 274	mrg_apply_transition
11.31	2.64	122 880	generate_4way_bernoulli
8.02	2.58	124 906	mrg_orig_step
2.49	2.49	173 454	mod_mul
1.89	1.89	72 822	mod_mac3
1.58	1.58	72 822	mod_mac2



當依照 self 來排序時，可以看到 verify\_bfs\_tree 有最高的 cost。

(2)



用 called 來排序，可以看到 mod\_mac 和 mod\_mac\_y 被呼叫最多次，mod\_mac 的 caller 是 mod\_mac\_y，mod\_mac\_y 的 caller 是 mrg\_orig\_step 和 mrg\_apply\_transition。

## 5. Pytorch profiler

(1)

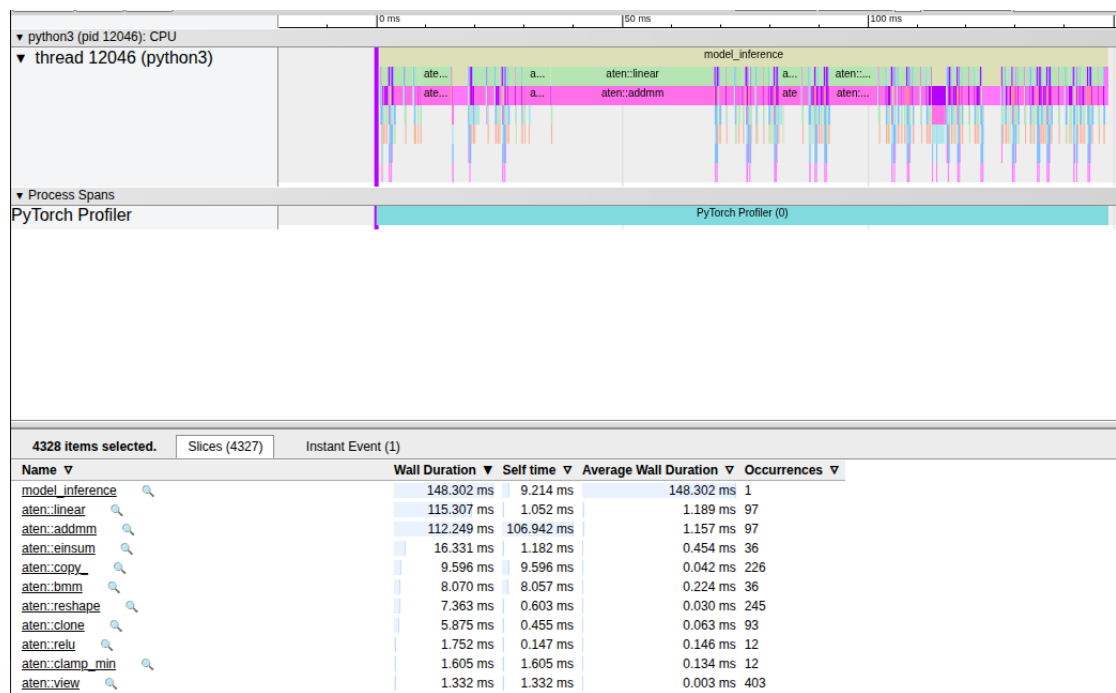
```
karrs@MSS: $ python3 test.py
cpu
torch.Size([2, 7, 10])
STAGE:2025-05-13 20:13:31 9172:9172 ActivityProfilerController.cpp:335] Completed Stage: Warm Up
STAGE:2025-05-13 20:13:31 9172:9172 ActivityProfilerController.cpp:342] Completed Stage: Collection
STAGE:2025-05-13 20:13:31 9172:9172 ActivityProfilerController.cpp:347] Completed Stage: Post Processing
```

Name	Self CPU %	Self CPU	CPU total %	CPU total	CPU time avg	CPU Mem	Self CPU Mem	# of Calls
aten::addmm	65.38%	60.530ms	67.26%	62.273ms	641.990us	4.13 Mb	4.13 Mb	97
aten::clamp_min	8.66%	8.021ms	8.66%	8.021ms	668.436us	1.50 Mb	1.50 Mb	12
model_inference	7.11%	6.582ms	100.00%	92.582ms	92.582ms	0 b	-10.72 Mb	1
aten::native_layer_norm	5.50%	5.088ms	5.71%	5.287ms	176.229us	939.66 Kb	72 b	30
aten::copy_	2.89%	2.676ms	2.89%	2.676ms	11.840us	0 b	0 b	226
aten::view	0.99%	917.073us	0.99%	917.073us	2.276us	0 b	0 b	403
aten::add	0.99%	916.528us	0.99%	916.528us	28.642us	1000.00 Kb	1000.00 Kb	32
aten::bmm	0.93%	864.130us	0.94%	873.837us	24.273us	624.38 Kb	624.38 Kb	36
aten::einsum	0.79%	734.221us	5.05%	4.673ms	129.815us	2.32 Mb	0 b	36
aten::linear	0.68%	632.436us	69.49%	64.331ms	663.203us	4.13 Mb	0 b	97

Self CPU time total: 92.582ms

前三名分別是 aten::addmm、aten::clamp\_min、aten::native\_layer\_norm，aten::addmm 是做矩陣乘法後加法；aten::clamp\_min 是設定一個 lower bound，將小於該 lower bound 的資料回傳 lower bound；aten::native\_layer\_norm 是在實作 layer normalization。

(2)



可以看到出現最多的是 `aten::linear` 和 `aten::addmm`，`aten::linear` 是用來實作全連接層；`aten::addmm` 是做矩陣乘法後加法。